

Kiran Kumar Mani

Dallas, Texas • +1-469-264-0196 • kiran_kumar_1991@yahoo.co.in

<https://www.linkedin.com/in/kirankumarmani> • https://bitbucket.org/kiran_kumar_1991/

EDUCATION

The University of Texas at Dallas

Masters in Computer Science. GPA: 3.83/4

Expected Graduation – May 2017

P.S.G College of Technology, Coimbatore, India

Bachelor of Engineering in Computer Science and Engineering. CGPA: 8.51/10

Graduation Year – May 2013

TECHNICAL EXPERTISE

Languages : Java, JavaScript, Node.js and Python

NoSQL : Cassandra and MongoDB

Cache : Memcached

RDBMS : MySQL and MariaDB Galera Cluster

Frameworks : Hibernate and SQLAlchemy

WORK EXPERIENCE

[24]7 iLabs

Jun 2013 – Jun 2015 – 2 years

Member of Technical Staff, Platform Engineering Team

- Worked on American Express SSO Integration with [24]7 Assist. A high priority client requirement and was successfully released to production within 3 months.
- Part of the 4-member core team which spearheaded the development of an in-house distributed data pipeline. The pipeline was built to replace Apache Kafka and has support for data quality checks, transformations, partitioning and streaming of data. The pipeline was written in Java and Cassandra was used to persist the data
- The pipeline supports horizontal scaling and guarantees zero data loss
- Achieved a throughput of 1300 events per second per node.
- Served as the backbone for suite of [24]7's products delivering near real-time performance

INTERNSHIPS

Juniper Networks

Jun 2016 – Aug 2016 – 3 months

Software Engineering Intern, Network Management Team

- Worked on Software Defined Secure Networks (SDSN) Policy Controller
- Evaluated the existing design of the feed connector component, which is used by the networking devices to fetch the threat feeds
- Replaced PostgreSQL with MariaDB Galera Cluster and used SQLAlchemy with connection pooling for interacting with the database cluster
- Redesigned the existing modules to support horizontal scaling and high availability

Microsoft Research India

Dec 2012 – May 2013 – 6 months

Research Intern, Advanced Development Group

- In my ramp up project, I worked on the Windows Azure implementation of an end to end data compression algorithm to handle large amount of data. We benchmarked the various communication methods between the Azure instances and logged the application stats to Azure Storage account.
- In my main project, along with my mentor, we used Machine Learning to predict the storage and the compute usage in the Azure data centers, which would help in capacity planning. The model had an accuracy of 95%. In addition to this, we analyzed the usage pattern of the Microsoft Azure customers during their trial period and predicted with 90% accuracy on whether the customer would buy a subscription or not.

ACHIEVEMENTS

- Microsoft Student Partner (MSP) for the academic years 2012 and 2013.
- Member of the Digital Literacy Program (DLP). Taught the under-privileged students about the basics of computers.
- Won first place for creating a video to spread awareness about Kidney Cancer